

Week 5: Data Representation

References: *Patterson & Hennessy, ARM Ed., Chapter 2*

Review

Last week, we studied the **building blocks of digital logic**: gates, combinational circuits such as multiplexers, decoders, and adders, and sequential circuits such as latches, flip-flops, and registers.

These components provided the hardware foundation for how a CPU performs arithmetic and controls the flow of data.

This week, we turn from **how** the circuits operate to **what** they operate on: data itself.

The 0s and 1s flowing through logic gates are not arbitrary.

They represent integers, signed and unsigned values, characters, and eventually floating-point numbers.

Understanding how data is encoded in binary is essential for bridging the gap between high-level programming concepts and the physical machine.

1. Bits, Bytes, and Words (§2.1 – §2.3)

The most fundamental unit of information in a computer is the **bit**, which can hold a value of 0 or 1.

Physically, this is implemented as a low or high voltage.

- A **byte** is a collection of 8 bits.
The byte is the smallest *addressable unit* of memory in most architectures. Even if only a single bit is required, the memory system allocates at least one byte.
- A **word** is the “natural” data size for a processor.
A word matches the width of the CPU registers and datapath.
 - In a 32-bit ARMv7 processor, a word is 32 bits (4 bytes).
 - In a 64-bit ARMv8 processor, a word is 64 bits (8 bytes).

For example, the decimal value 300_{10} is equal to 100101100_2 .

This requires 9 bits to store.

Since memory can only address whole bytes, at least 2 bytes are required.

Memory alignment

Processors often require that words be stored at memory addresses that are multiples of the word size.

For example, a 4-byte word should be placed at an address divisible by 4. Misaligned accesses may be slower or may even cause exceptions on some systems. Alignment ensures efficient memory access.

2. Unsigned Binary Numbers (§2.4)

Binary numbers are represented using **positional notation**, similar to decimal numbers but with base 2.

Each position corresponds to a power of 2.

For an n -bit binary number:

$$\text{Value} = (b_{n-1} \times 2^{n-1}) + (b_{n-2} \times 2^{n-2}) + \dots + (b_1 \times 2^1) + (b_0 \times 2^0)$$

Example:

$$1101_2 = (1 \times 8) + (1 \times 4) + (0 \times 2) + (1 \times 1) = 13_{10}$$

Range of unsigned values

- An n -bit unsigned integer can represent values from 0 to $2^n - 1$.
- Examples:
 - 8-bit unsigned: $0 \rightarrow 255$
 - 16-bit unsigned: $0 \rightarrow 65,535$
 - 32-bit unsigned: $0 \rightarrow 4,294,967,295$

Overflow in unsigned arithmetic

Since registers have fixed width, results that exceed the maximum representable value **wrap around**.

Example (8-bit):

$$11111111_2 (255) + 1 = 00000000_2 (0)$$

The 9th bit is discarded, producing wraparound.

In C and most modern languages, unsigned arithmetic is explicitly defined as modulo 2^n .

3. Signed Numbers (§2.4)

Unsigned binary cannot represent negative values.

Several historical methods were developed to represent signed integers, but modern systems universally use **two's complement**.

3.1 Sign-and-Magnitude

The most intuitive scheme uses the leftmost bit to indicate sign (0 = positive, 1 = negative), with the remaining bits giving the magnitude.

- Example (4-bit): $1001_2 = -9$.
- Problem: Two encodings for zero ($0000 = +0$, $1000 = -0$).

3.2 One's Complement

Negative numbers are obtained by flipping all the bits of the positive representation.

- Example (4-bit): $+5 = 0101$, $-5 = 1010$.
- Problem: Still has two encodings for zero.

3.3 Two's Complement (standard today)

Negative numbers are represented by inverting all bits of the positive number and then adding 1.

- Example (4-bit):
 - $+6 = 0110$
 - $-6 = \text{invert}(0110) = 1001 + 1 = 1010$

Range of two's complement

An n -bit two's complement number represents values from -2^{n-1} to $2^{n-1} - 1$.

- 4-bit: $-8 \rightarrow +7$
- 32-bit: $-2,147,483,648 \rightarrow +2,147,483,647$

Why two's complement?

- Only one representation for zero.
- Addition and subtraction use the **same adder circuits** as unsigned numbers.
- The carry out from the most significant bit can be ignored.

Code Snippet

```
#include <stdint.h>
#include <stdio.h>

int8_t negate8(int8_t x){
    return (int8_t)((~x) + 1);    // two's complement negation in 8 bits
}

int main(void){
    int8_t a = 53;                // 0b0011_0101
    int8_t na = negate8(a);       // should be -53
    printf("a=%d neg(a)=%d\n", a, na);
    return 0;
}
```

4. Arithmetic in Two's Complement (§2.4)

Two's complement arithmetic follows the same binary addition rules, with results interpreted as signed or unsigned depending on context.

Example of overflow

$$0111_2 (7) + 0001_2 (1) = 1000_2 (-8)$$

The 4-bit signed range is $-8 \rightarrow +7$, so the result has overflowed.

Detecting overflow

- If two positive numbers yield a negative result, overflow has occurred.
- If two negative numbers yield a positive result, overflow has occurred.
- If the operands have different signs, overflow cannot occur.

Code Snippet

```
#include <stdint.h>
#include <stdio.h>

int add_overflows(int32_t a, int32_t b){
    int32_t s = a + b;
    return ((a ^ s) & (b ^ s)) < 0;    // same signs in, different sign out
}
```

```

int main(void){
    int32_t x = 1<<30;    // large positive
    int32_t y = 1<<30;    // large positive
    int32_t s = x + y;
    printf("s=%d overflow=%d\n", s, add_overflows(x, y)); // expect overflow=1
    return 0;
}

```

Subtraction

Subtraction is performed by adding the two's complement of the subtrahend.

Example: $5 - 3$ is equivalent to $5 + (-3)$.

Code Snippet

```

#include <stdint.h>
#include <stdio.h>

int32_t sub_via_add(int32_t a, int32_t b){
    return a + (~b + 1); // two's complement subtraction
}

int main(void){
    int32_t a = 5, b = 3;
    printf("a-b=%d sub_via_add=%d\n", a-b, sub_via_add(a,b));
    return 0;
}

```

5. Endianness (§2.6, §2.8 supplement)

Endianness refers to the ordering of bytes when multi-byte data is stored in memory.

- **Little-endian:** least significant byte stored at the lowest memory address.
This is the default on ARM and x86.
- **Big-endian:** most significant byte stored at the lowest address.

Example: Storing the 32-bit word 0x12345678 at address 0x1000.

- Little-endian:
 - 0x1000: 78

- 0x1001: 56
- 0x1002: 34
- 0x1003: 12
- Big-endian:
 - 0x1000: 12
 - 0x1001: 34
 - 0x1002: 56
 - 0x1003: 78

Endianness is important for interpreting memory dumps, file formats, and network protocols.

6. Characters and ASCII (§2.9)

Computers must also represent non-numeric data such as text. This is achieved by mapping characters to integers.

- **ASCII (American Standard Code for Information Interchange):** A 7-bit encoding for 128 characters, including letters, digits, punctuation, and control codes.
 - Examples: 'A' = 65, 'a' = 97, newline = 10.
- **Extended ASCII:** Uses 8 bits to allow 256 characters. Different systems defined different extended sets.
- **Unicode:** Defines a much larger set of characters for all writing systems.
- **UTF-8:** A variable-length encoding for Unicode. Backward compatible with ASCII: the first 128 characters have identical byte values.

Transition to Week 6

This week we established how **integers and characters** are represented inside a computer.

We examined unsigned binary numbers, extended the system to negatives with two's complement, and explored how overflow arises in fixed-width arithmetic. We also saw how characters are encoded in ASCII and Unicode, connecting binary storage to human-readable text.

Next week, we expand this foundation in two directions.

First, we study **bitwise operations**, which treat numbers as raw bit patterns and allow direct manipulation of individual bits — a key tool for systems programming and hardware design.

Second, we turn to **floating-point representation**, which extends binary arithmetic to fractions and real numbers using the IEEE 754 standard.

Together, these topics prepare us to handle both **low-level bit manipulation** and the **wide range of values** required in real-world scientific and engineering applications.